

ENV 790.30 - Time Series Analysis for Energy Data | Spring 2023

Assignment 8 - Due date 03/27/23

Qiuying Liao

Directions

You should open the .rmd file corresponding to this assignment on RStudio. The file is available on our class repository on Github. And to do so you will need to fork our repository and link it to your RStudio.

Once you have the project open the first thing you will do is change “Student Name” on line 3 with your name. Then you will start working through the assignment by **creating code and output** that answer each question. Be sure to use this assignment document. Your report should contain the answer to each question and any plots/tables you obtained (when applicable).

When you have completed the assignment, **Knit** the text and code into a single PDF file. Rename the pdf file such that it includes your first and last name (e.g., “LuanaLima_TSA_A08_Sp22.Rmd”). Submit this pdf using Sakai.

Set up

Some packages needed for this assignment: `forecast`, `tseries`, `smooth`. Do not forget to load them before running your script, since they are NOT default packages.

```
#Load/install required package here
```

```
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method             from
```

```
##   as.zoo.data.frame zoo
```

```
library(tseries)
```

```
library(smooth)
```

```
## Loading required package: greybox
```

```
## Package "greybox", v1.0.7 loaded.
```

```
## This is package "smooth", v3.2.0
```

```
library(tidyr)
```

```
##
```

```
## Attaching package: 'tidyr'
```

```
## The following object is masked from 'package:greybox':
```

```
##
```

```
##   spread
```

```
library(lubridate)
```

```
##
```

```
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:greybox':
##
##      hm
## The following objects are masked from 'package:base':
##
##      date, intersect, setdiff, union
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##      filter, lag
## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
library(ggplot2)
library(kableExtra)

## Warning in !is.null(rmarkdown::metadata$output) && rmarkdown::metadata$output
## %in% : 'length(x) = 2 > 1' in coercion to 'logical(1)'
##
## Attaching package: 'kableExtra'
## The following object is masked from 'package:dplyr':
##
##      group_rows
```

Importing and processing the data set

Consider the data from the file “inflowtimeseries.txt”. The data corresponds to the monthly inflow in m^3/s for some hydro power plants in Brazil. You will only use the last column of the data set which represents one hydro plant in the Amazon river basin. The data span the period from January 1931 to August 2011 and is provided by the Brazilian ISO.

For all parts of the assignment prepare the data set such that the model consider only the data from January 2000 up to December 2009. Leave the year 2010 of data (January 2010 to December 2010) for the out-of-sample analysis. Do **NOT** use data from 2010 and 2011 for model fitting. You will only use it to compute forecast accuracy of your model.

Part I: Preparing the data sets

Q1

Read the file into a data frame. Prepare your time series data vector such that observations start in January 2000 and end in December 2009. Make you sure you specify the **start=** and **frequency=** arguments. Plot the time series over time, ACF and PACF.

```
inflow_data <- read.table(file="../Data/inflowtimeseries.txt", header=FALSE, skip=0)

#last column of the table and date
inflow_data1 <- inflow_data[,1:2]
inflow_data2 <- inflow_data[,17]
inflow_data_hydro <- cbind(inflow_data1, inflow_data2)
```

```

#add column names
colnames(inflow_data_hydro)=c("Month", "Year", "HP17")

#Date
inflow_data_hydro$Month <- unite(inflow_data_hydro, col = 'Time', c('Month','Year'), sep=" ")
inflow_data_hydro <- inflow_data_hydro[,1:1]
inflow_data_hydro$Time <- my(inflow_data_hydro$Time)

#January 2000 up to December 2009
inflow_data_hydro_2009 <- inflow_data_hydro[829:948,]
head(inflow_data_hydro_2009)

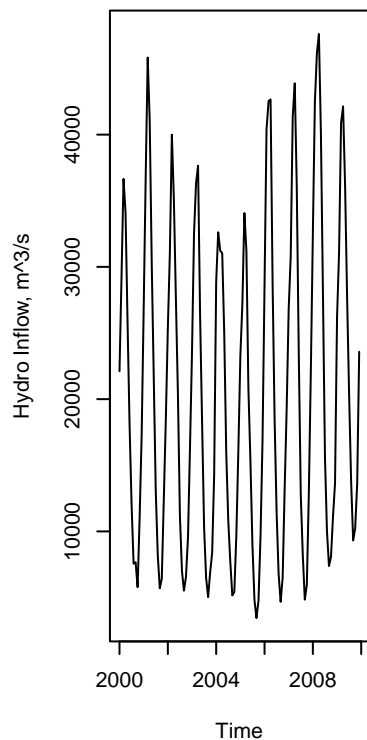
##           Time  HP17
## 829 2000-01-01 22107
## 830 2000-02-01 29247
## 831 2000-03-01 36651
## 832 2000-04-01 34089
## 833 2000-05-01 25821
## 834 2000-06-01 18007

#ts
ts_inflow_data_hydro_2009 <- ts(inflow_data_hydro_2009[,2], start = c(2000,01), frequency = 12)


#Acf and Pacf
par(mfrow=c(1,3))
plot(ts_inflow_data_hydro_2009, xlab = "Time", ylab = "Hydro Inflow, m^3/s", main = "Hydro Inflow of a power plant from 2000 to 2009")
Acf(ts_inflow_data_hydro_2009, lag =40, main = "Acf of hydro inflow")
Pacf(ts_inflow_data_hydro_2009, lag =40, main = "Pacf of hydro inflow")

```


Hydro Inflow of a power plant from 2000 to 2009



Acf of hydro inflow



Pacf of hydro inflow




Q2

Using the `decompose()` or `stl()` and the `seasadj()` functions create a series without the seasonal component, i.e., a deseasonalized inflow series. Plot the deseasonalized series and original series together using `ggplot`, make sure your plot includes a legend. Plot ACF and PACF for the deseasonalized series. Compare with the plots obtained in Q1.

```
#decompose
decompose_hydro <- decompose(ts_inflow_data_hydro_2009, "additive")
plot(decompose_hydro)
```

Decomposition of additive time series



```
print(decompose_hydro)
```

```
## $x
##      Jan   Feb   Mar   Apr   May   Jun   Jul   Aug   Sep   Oct   Nov   Dec
## 2000 22107 29247 36651 34089 25821 18007 12027  7557  7670  5777 11116 17269
## 2001 26835 36764 45828 41038 30282 21745 13040  8204  5697  6429 12473 18447
## 2002 25356 31078 39996 35027 27575 20372 11153  7020  5519  6563  9518 16380
## 2003 23553 32547 36319 37651 26376 18267 10274  6488  5029  6888  8393 14285
## 2004 28901 32620 31225 31042 24951 15966 10793  7711  5159  5424  9656 16103
## 2005 23242 27204 34070 31232 20811 15307  8992  4820  3449  4742  9516 16836
## 2006 28607 40407 42531 42658 28826 18517 11419  6766  4675  6499 13263 19300
## 2007 26840 30748 41345 43887 36150 23392 12742  8168  4838  5956 12880 22223
## 2008 32493 42492 46138 47613 40093 28051 15616 10012  7378  8083 11131 13617
## 2009 26184 31029 40892 42138 36104 27713 20302 13772  9316 10225 13586 23583
##
## $seasonal
##      Jan      Feb      Mar      Apr      May
## 2000  5992.19367 12911.55478 18814.60571 18113.28164  9068.14275
## 2001  5992.19367 12911.55478 18814.60571 18113.28164  9068.14275
## 2002  5992.19367 12911.55478 18814.60571 18113.28164  9068.14275
```

##	2003	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##	2004	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##	2005	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##	2006	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##	2007	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##	2008	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##	2009	5992.19367	12911.55478	18814.60571	18113.28164	9068.14275			
##		Jun	Jul	Aug	Sep	Oct			
##	2000	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2001	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2002	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2003	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2004	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2005	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2006	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2007	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2008	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##	2009	-65.63503	-8741.35262	-13136.25540	-15089.91744	-14374.92670			
##		Nov	Dec						
##	2000	-9839.24151	-3652.44985						
##	2001	-9839.24151	-3652.44985						
##	2002	-9839.24151	-3652.44985						
##	2003	-9839.24151	-3652.44985						
##	2004	-9839.24151	-3652.44985						
##	2005	-9839.24151	-3652.44985						
##	2006	-9839.24151	-3652.44985						
##	2007	-9839.24151	-3652.44985						
##	2008	-9839.24151	-3652.44985						
##	2009	-9839.24151	-3652.44985						
##									
##	\$trend								
##		Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
##	2000	NA	NA	NA	NA	NA	NA	19141.83	19652.04
##	2001	22034.54	22103.71	22048.46	21993.42	22077.12	22182.75	22170.21	21871.67
##	2002	20229.21	20101.25	20044.50	20042.67	19925.12	19715.87	19554.62	19540.71
##	2003	19202.21	19143.42	19100.83	19093.96	19060.62	18926.46	19062.00	19287.88
##	2004	18026.79	18099.37	18155.75	18100.17	18091.79	18220.17	18060.12	17598.67
##	2005	17150.96	16955.46	16763.75	16664.08	16629.83	16654.54	16908.62	17682.29
##	2006	20926.21	21108.42	21240.58	21364.88	21594.21	21853.00	21882.04	21405.96
##	2007	22078.79	22192.33	22257.54	22241.71	22203.12	22308.96	22666.29	23391.17
##	2008	25427.00	25623.58	25806.25	26000.71	26016.46	25585.00	24963.54	24223.04
##	2009	22686.67	23038.58	23276.00	23446.00	23637.54	24155.08	NA	NA
##		Sep	Oct	Nov	Dec				
##	2000	20347.62	21019.54	21494.96	21836.58				
##	2001	21391.75	20898.29	20535.04	20365.04				
##	2002	19448.71	19404.83	19464.21	19326.54				
##	2003	19078.67	18591.04	18256.29	18101.04				
##	2004	17491.54	17618.00	17453.42	17253.46				
##	2005	18584.96	19413.58	20223.62	20691.33				
##	2006	20954.08	20955.88	21312.25	21820.54				
##	2007	24080.21	24435.17	24754.71	25113.12				
##	2008	23526.83	23080.12	22685.79	22505.50				
##	2009	NA	NA	NA	NA				
##									

```


## $random
##           Jan           Feb           Mar           Apr           May           Jun
## 2000         NA         NA         NA         NA         NA         NA
## 2001 -1191.73534  1748.73688  4964.93596   931.30170  -863.26775  -372.11497
## 2002  -865.40201 -1934.80478  1136.89429 -3128.94830 -1418.26775   721.76003
## 2003 -1641.40201   492.02855 -1596.43904   443.76003 -1752.76775  -593.82330
## 2004  4882.01466  1609.07022 -5745.35571 -5171.44830 -2208.93441 -2188.53164
## 2005   98.84799 -2663.01312 -1508.35571 -3545.36497 -4886.97608 -1281.90664
## 2006  1688.59799  6387.02855  2475.81096  3179.84336 -1836.35108 -3270.36497
## 2007 -1230.98534 -4355.88812   272.85262  3532.01003  4878.73225  1148.67670
## 2008  1073.80633  3956.86188  1517.14429  3499.01003  5008.39892  2531.63503
## 2009 -2494.86034 -4921.13812 -1198.60571   578.71836  3398.31559  3623.55170
##           Jul           Aug           Sep           Oct           Nov           Dec
## 2000  1626.51929  1041.21373  2412.29244  -867.61497  -539.71682  -915.13349
## 2001  -388.85571  -531.41127  -604.83256   -94.36497  1777.19985  1734.40818
## 2002   339.72762   615.54707  1160.20910  1533.09336  -106.96682   705.90818
## 2003  -46.64738   336.38040  1040.25077  2671.88503  -24.05015  -163.59182
## 2004  1474.22762  3248.58873  2757.37577  2180.92670  2041.82485  2501.99151
## 2005   824.72762   273.96373   -46.04090  -296.65664  -868.38349  -202.88349
## 2006 -1721.68904 -1503.70293 -1189.16590   -81.94830  1789.99151  1131.90818
## 2007 -1182.93904 -2086.91127 -4152.29090 -4104.23997 -2035.46682   762.32485
## 2008  -606.18904 -1074.78627 -1058.91590  -622.19830 -1715.55015 -5236.05015
## 2009         NA         NA         NA         NA         NA         NA
##
## $figure
## [1] 5992.19367 12911.55478 18814.60571 18113.28164 9068.14275
## [6] -65.63503 -8741.35262 -13136.25540 -15089.91744 -14374.92670
## [11] -9839.24151 -3652.44985
##
## $type
## [1] "additive"
##
## attr("class")
## [1] "decomposed.ts"

```

```

#create non-seasonal natural gas time series
#seasonal adjustment to original data, only works for decompose
deseasonal_hydro <- seasadj(decompose_hydro)
plot(deseasonal_hydro)

```




#Plot the deseasonalized series and original series together using ggplot with legend

```
ggplot(data = inflow_data_hydro_2009, aes(x=Time, y=inflow_data_hydro_2009[,2]))+
  geom_line(aes(x=Time, y=ts_inflow_data_hydro_2009, color = "Original"))+
  geom_line(aes(x=Time, y=deseasonal_hydro, color = "Deseasonal"))+
  ylab("Hydro Inflow, m3/s")+
  xlab("Year")+
  ggtitle("Monthly inflow of a hydro power plant in the Amazon River basin")+
  scale_color_manual(name = "Types", values = c("Original" = "blue", "Deseasonal" = "red"), labels=c("Original", "Deseasonal"))
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```


Monthly inflow of a hydro power plant in the Amazon River basin



```
#Comparing ACFs  
#seasonality is removed  
par(mar=c(3,3,3,0));par(mfrow=c(1,2))  
Acf(deseasonal_hydro,lag.max=40,main="ACF of Non Sesonal Hydro Flow")  
Pacf(deseasonal_hydro,lag.max=40,main="PACF of Non Sesonal Hydro Flow")
```


ACF of Non Seasonal Hydro Flo

PACF of Non Seasonal Hydro Flc



Answer: Compared the plots in Q1, the acf of nonseasonal hydro flow have slow decay instead of seasonal spike. Therefore, we removed the seasonality. Also, we removed those little spikes for the seasonality in the pacf of nonseasonal hydro flow compared to the pacf of hydro flow.

Part II: Forecasting with ARIMA models and its variations

Q3

Fit a non-seasonal ARIMA(p, d, q) model using the `auto.arima()` function to the non-seasonal data. Forecast 12 months ahead of time using the `forecast()` function. Plot your forecasting results and further include on the plot the last year of non-seasonal data to compare with forecasted values (similar to the plot on the lesson file for M10).

```
#auto arima
ARIMA_autofit <- auto.arima(deseasonal_hydro,max.D=0,max.P = 0,max.Q=0)
print(ARIMA_autofit)
```

```
## Series: deseasonal_hydro
## ARIMA(0,1,0)
##
## sigma^2 = 5031380: log likelihood = -1087.01
## AIC=2176.02 AICc=2176.05 BIC=2178.8
```

```
#forecast
ARIMA_forecast <- forecast(object = ARIMA_autofit, h = 12)
print(ARIMA_forecast)
```

```
##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## Jan 2010      27235.45 24360.84 30110.06 22839.11 31631.79
## Feb 2010      27235.45 23170.13 31300.77 21018.08 33452.82
```

## Mar 2010	27235.45	22256.47	32214.43	19620.76	34850.14
## Apr 2010	27235.45	21486.22	32984.68	18442.76	36028.14
## May 2010	27235.45	20807.62	33663.28	17404.93	37065.97
## Jun 2010	27235.45	20194.11	34276.79	16466.65	38004.25
## Jul 2010	27235.45	19629.93	34840.97	15603.82	38867.08
## Aug 2010	27235.45	19104.81	35366.09	14800.71	39670.19
## Sep 2010	27235.45	18611.61	35859.29	14046.42	40424.48
## Oct 2010	27235.45	18145.12	36325.78	13332.99	41137.91
## Nov 2010	27235.45	17701.43	36769.47	12654.43	41816.47
## Dec 2010	27235.45	17277.49	37193.41	12006.07	42464.83

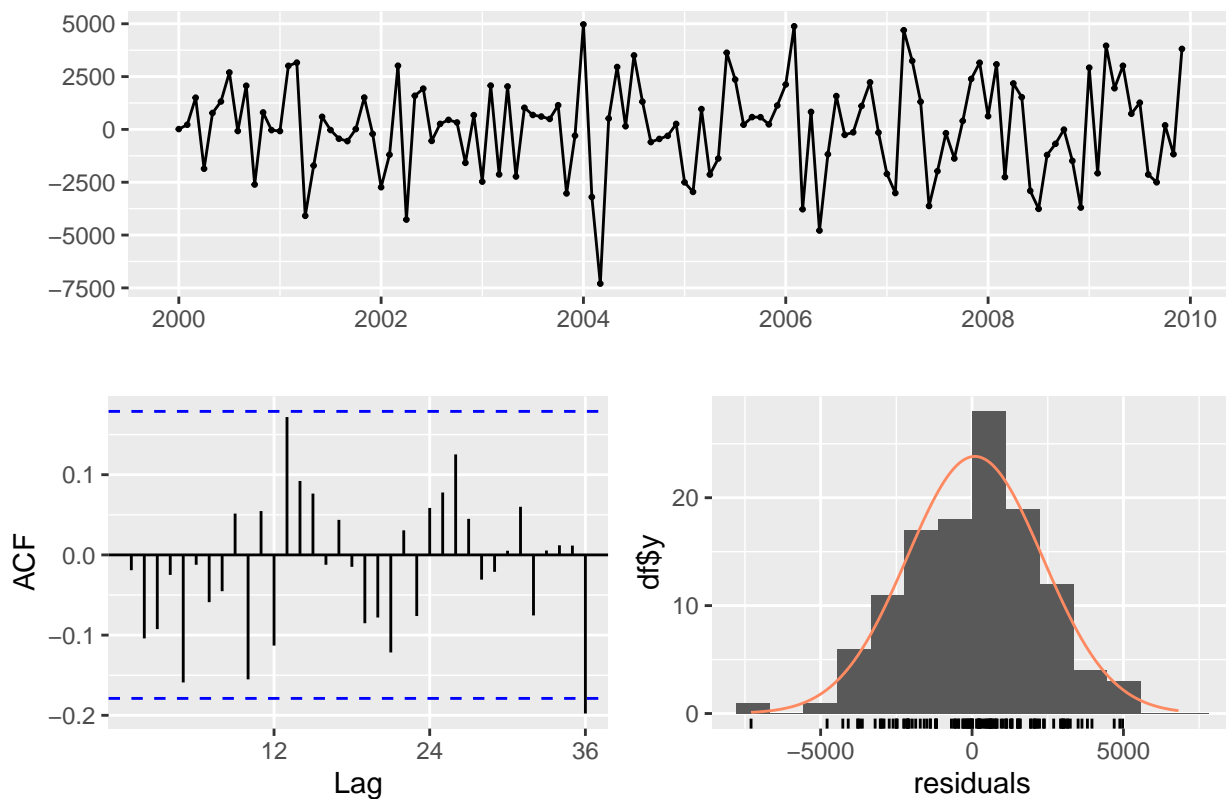
```
plot(ARIMA_forecast)
```

Forecasts from ARIMA(0,1,0)



```
checkresiduals(ARIMA_forecast)
```

Residuals from ARIMA(0,1,0)




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,0)
## Q* = 24.208, df = 24, p-value = 0.4498
##
## Model df: 0.   Total lags used: 24

#plot your forecasting results and further include on the plot the last year of non-seasonal data to compare
#create time series of 2010 data
#January 2000 up to December 2010
inflow_data_hydro_2010 <- inflow_data_hydro[829:960,]
#ts
ts_inflow_data_hydro_2010 <- ts(inflow_data_hydro_2010[,2], start = c(2000,01), frequency = 12)
#decompose
decompose_hydro_2010 <- decompose(ts_inflow_data_hydro_2010, "additive")
#seasonal adjustment to original data, only works for decompose
deseasonal_hydro_2010 <- seasadj(decompose_hydro_2010)
#only for 2010
deseasonal2_hydro_2010 <- window(deseasonal_hydro_2010, start = c(2010,01), end = c(2010,12))

autoplot(deseasonal_hydro, series="Deseasonal 2000 to 2009",PI=FALSE) +
  autolayer(deseasonal2_hydro_2010,series="Deseasonal for 2010",PI=FALSE) +
  autolayer(ARIMA_forecast,series="Forecast",PI=FALSE) +
  ylab("Hydro flow, m^3/s") +
  xlab("Year") +
  labs(col="Models")
```

```
## Warning in ggplot2::geom_line(ggplot2::aes_(group = ~series, colour =
## ~series), : Ignoring unknown parameters: `PI`

## Warning in ggplot2::geom_line(ggplot2::aes_(x = ~timeVal, y = ~seriesVal, :
## Ignoring unknown parameters: `PI`
```



Q4

Put the seasonality back on your forecasted values and compare with the original seasonal data values. *Hint* : One way to do it is by summing the last year of the seasonal component from your decompose object to the forecasted series.

```
#Seasonal component from decompose
Seasonal_decompose_inflow <- as.data.frame(decompose_hydro$seasonal)
Seasonal_decompose_inflow1 <- cbind(inflow_data_hydro_2009[,1], Seasonal_decompose_inflow)
Sea_decom_inflow_2009 <- Seasonal_decompose_inflow[109:120,]

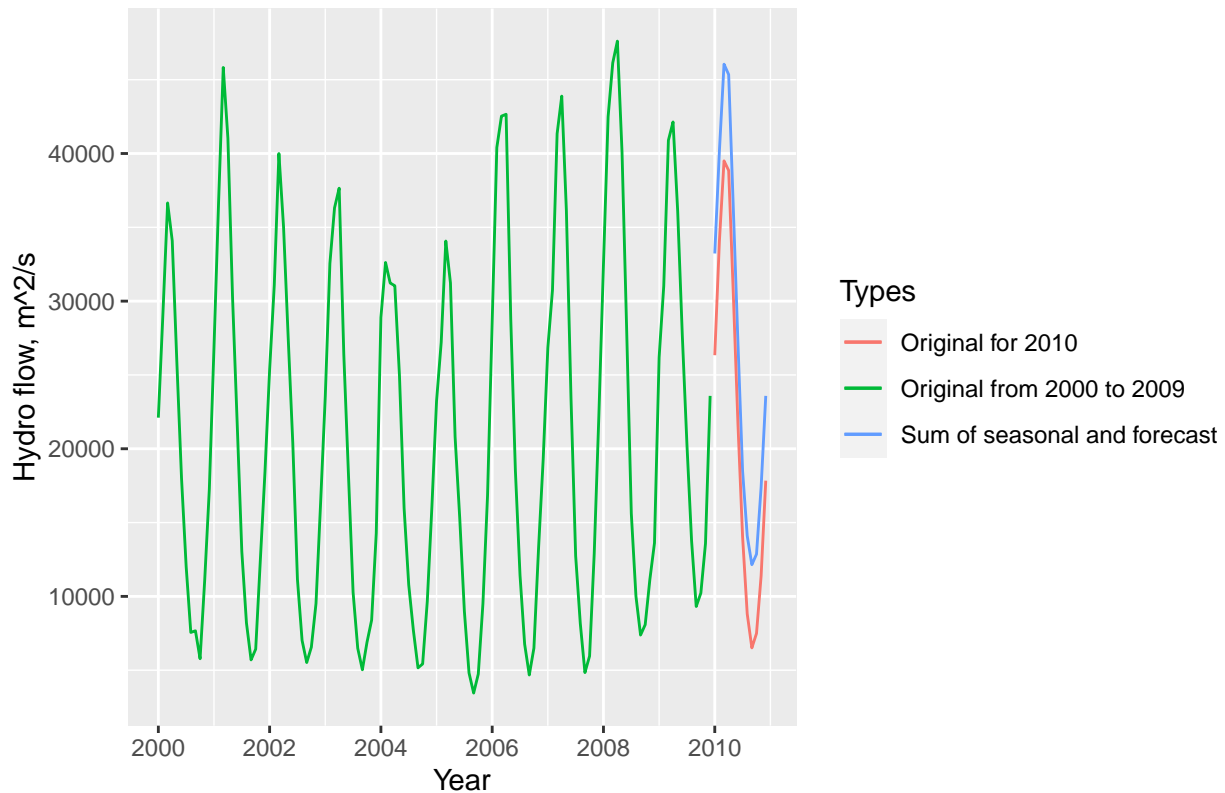
#Add seasonality to the forecasted series
Sum_Sea_forecast <- Sea_decom_inflow_2009 + ARIMA_forecast$mean

#original seasonal data values (2010)
inflow_data_only_2010 <- inflow_data_hydro[949:960,]
ts_inflow_data_only_2010 <- ts(inflow_data_only_2010[,2], start = c(2010,1), end = c(2010,12), frequency=12)

#plot
autoplot(ts_inflow_data_hydro_2009, series = "Original from 2000 to 2009") +
  autolayer(Sum_Sea_forecast, series="Sum of seasonal and forecast", PI=FALSE) +
  autolayer(ts_inflow_data_only_2010, series="Original for 2010", PI=FALSE) +
  ylab("Hydro flow, m³/s") +
  xlab("Year") +
```

```
labs(col="Types")
```

```
## Warning in ggplot2::geom_line(ggplot2::aes_(x = ~timeVal, y = ~seriesVal, : Ignoring unknown parameter
## Ignoring unknown parameters: `PI`
```



Q5

Repeat Q3 for the original data, but now fit a seasonal ARIMA(p, d, q)(P, D, Q)₁₂ also using the `auto.arima()`.

```
#auto sarima
SARIMA_autofit <- auto.arima(ts_inflow_data_hydro_2009)
print(SARIMA_autofit)

## Series: ts_inflow_data_hydro_2009
## ARIMA(1,0,0)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1      sma1      drift
##          0.7739  -0.8724  54.7963
## s.e.    0.0614   0.1767  25.8402
##
## sigma^2 = 5566545: log likelihood = -999.07
## AIC=2006.14   AICc=2006.52   BIC=2016.86


#forecast
SARIMA_forecast <- forecast(object = SARIMA_autofit, h = 12)
print(SARIMA_forecast)

##          Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Jan 2010          32655.388 29601.817 35708.96 27985.355 37325.42
```

```
## Feb 2010      38975.665 35118.619 42832.71 33076.822 44874.51
## Mar 2010      44694.285 40427.930 48960.64 38169.459 51219.11
## Apr 2010      43882.582 39388.955 48376.21 37010.173 50754.99
## May 2010      34835.315 30210.978 39459.65 27763.001 41907.63
## Jun 2010      25402.604 20701.808 30103.40 18213.357 32591.85
## Jul 2010      16934.482 12188.618 21680.35  9676.308 24192.66
## Aug 2010      12051.454  7279.018 16823.89  4752.643 19350.26
## Sep 2010       9546.441  4758.520 14334.36  2223.948 16868.93
## Oct 2010      10305.013  5508.438 15101.59  2969.285 17640.74
## Nov 2010      14706.158  9905.401 19506.92  7364.033 22048.28
## Dec 2010      21396.427 16594.823 26198.03 14053.008 28739.84
```

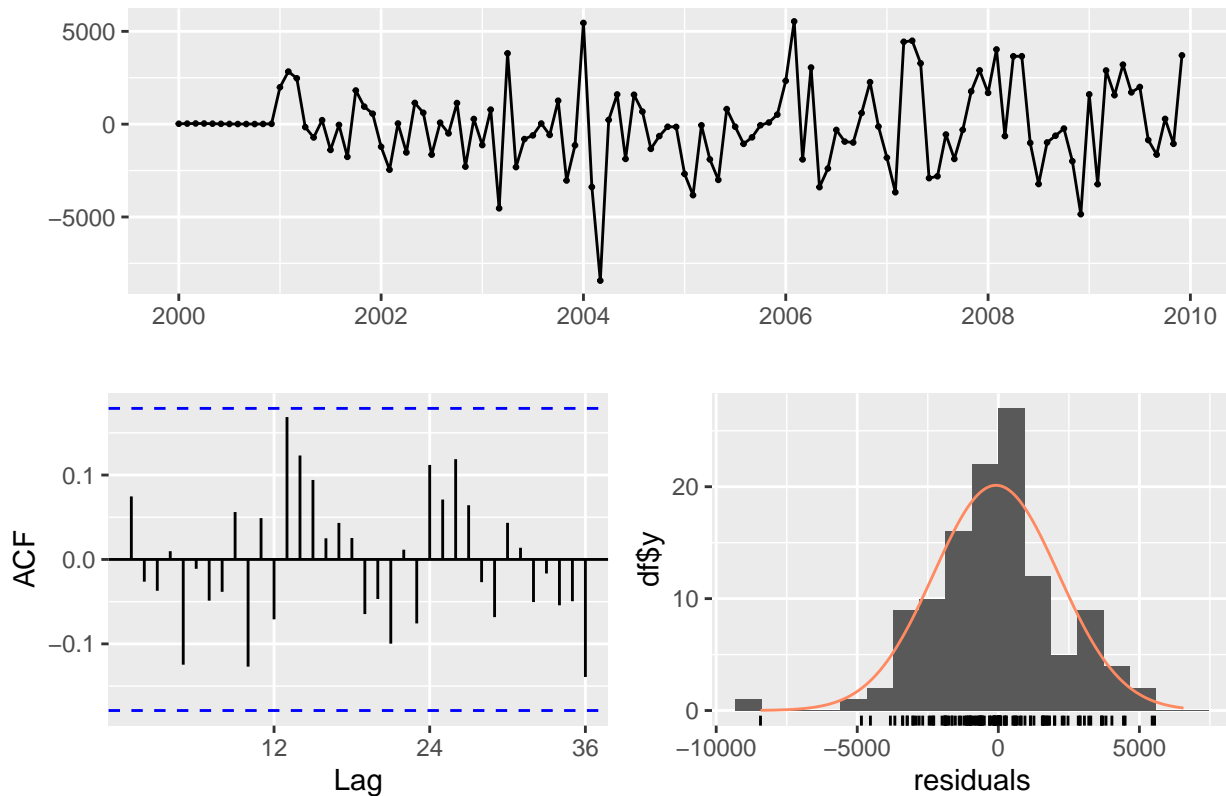
```
plot(SARIMA_forecast)
```

Forecasts from ARIMA(1,0,0)(0,1,1)[12] with drift



```
checkresiduals(SARIMA_forecast)
```

Residuals from ARIMA(1,0,0)(0,1,1)[12] with drift




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,0,0)(0,1,1)[12] with drift
## Q* = 19.881, df = 22, p-value = 0.5905
##
## Model df: 2.   Total lags used: 24
#plot your forecasting results and further include on the plot the last year of seasonal data to compar

autoplot(ts_inflow_data_hydro_2009, series="Original 2000 to 2009",PI=FALSE) +
  autolayer(ts_inflow_data_only_2010,series="Original Seasonal for 2010",PI=FALSE) +
  autolayer(SARIMA_forecast,series="Forecast",PI=FALSE) +
  ylab("Hydro flow, m^3/s") +
  xlab("Year") +
  labs(col="Models")

## Warning in ggplot2::geom_line(ggplot2::aes_(group = ~series, colour =
## ~series), : Ignoring unknown parameters: `PI`

## Warning in ggplot2::geom_line(ggplot2::aes_(x = ~timeVal, y = ~seriesVal, :
## Ignoring unknown parameters: `PI`
```

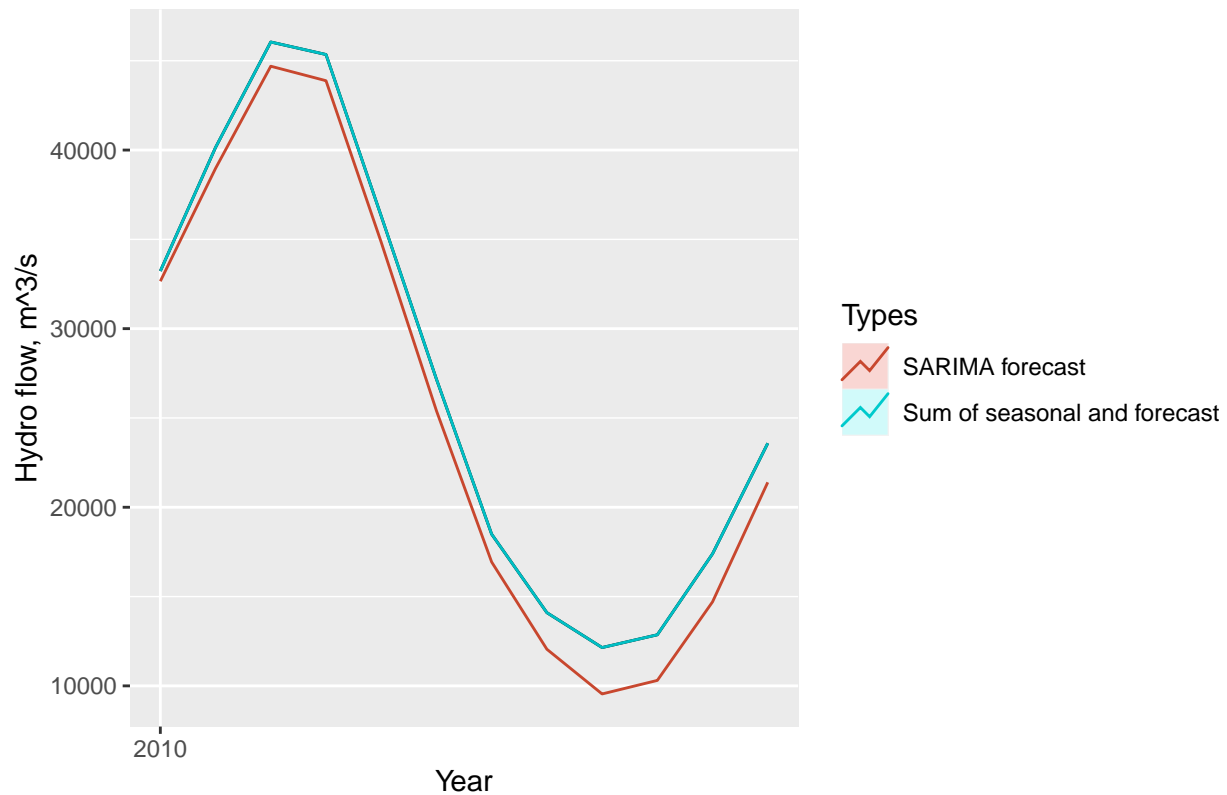


Q6

Compare the plots from Q4 and Q5 using the `autoplot()` function.

```
#plot
autoplot(Sum_Sea_forecast) +
  autolayer(Sum_Sea_forecast, series="Sum of seasonal and forecast", PI=FALSE) +
  autolayer(SARIMA_forecast, series="SARIMA forecast", PI=FALSE) +
  ylab("Hydro flow, m^3/s") +
  xlab("Year") +
  labs(col="Types")
```

```
## Warning in ggplot2::geom_line(ggplot2::aes_(x = ~timeVal, y = ~seriesVal, :
## Ignoring unknown parameters: `PI`
```

Part III: Forecasting with Other Models

Q7


Fit an exponential smooth model to the original time series using the function `ses()` from package `forecast`. Note that this function automatically do the forecast. Do not forget to set the arguments: `silent=FALSE` and `holdout=FALSE`, so that the plot is produced and the forecast is for the year of 2010.

```
# Exponential smooth model to the original ts
SES_original=ses(y = ts_inflow_data_hydro_2009, h = 12, holdout = FALSE, silent = FALSE)
print(SES_original)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2010	23582	14849.8153	32314.19	10227.276	36936.72
## Feb 2010	23582	11233.4433	35930.56	4696.512	42467.49
## Mar 2010	23582	8458.4205	38705.58	452.481	46711.52
## Apr 2010	23582	6118.9402	41045.06	-3125.445	50289.45
## May 2010	23582	4057.8032	43106.20	-6277.682	53441.68
## Jun 2010	23582	2194.3852	44969.62	-9127.534	56291.53
## Jul 2010	23582	480.7907	46683.21	-11748.251	58912.25
## Aug 2010	23582	-1114.1874	48278.19	-14187.559	61351.56
## Sep 2010	23582	-2612.2260	49776.23	-16478.612	63642.61
## Oct 2010	23582	-4029.1079	51193.11	-18645.546	65809.55
## Nov 2010	23582	-5376.7480	52540.75	-20706.583	67870.58
## Dec 2010	23582	-6664.4029	53828.40	-22675.882	69839.88

```
plot(SSES_original)
```

Forecasts from Simple exponential smoothing



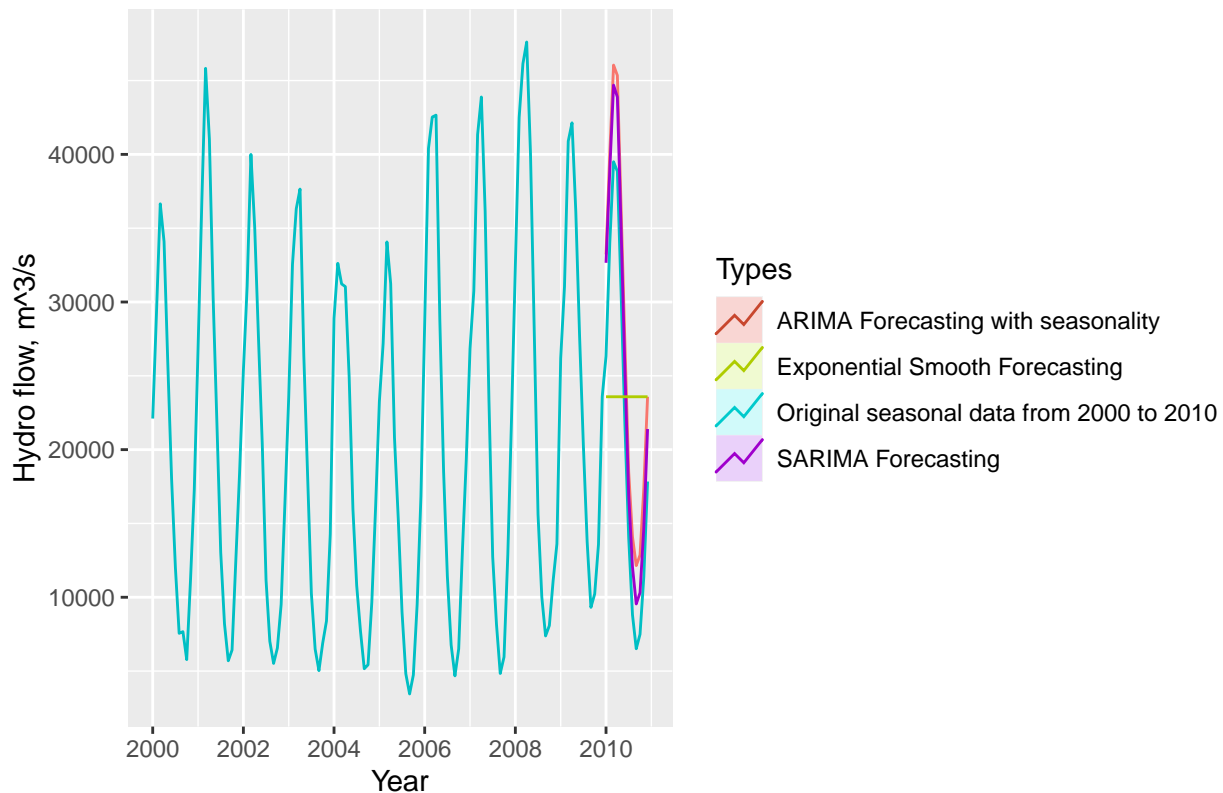
Part IV: Checking Forecast Accuracy

Q8

Make one plot with the complete original seasonal historical data (Jan 2000 to Dec 2010). Now add the forecasts from each of the developed models in parts Q4, Q5, Q7 and Q8. You can do it using the `autoplot()` combined with `autolayer()`. If everything is correct in terms of time line, the forecasted lines should appear only in the final year. If you decide to use `ggplot()` you will need to create a data frame with all the series will need to plot. Remember to use a different color for each model and add a legend in the end to tell which forecast lines corresponds to each model.

```
autoplot(ts_inflow_data_hydro_2010, series = "Original seasonal data from 2000 to 2010", PI=FALSE) +  
  autolayer(Sum_Sea_forecast, series="ARIMA Forecasting with seasonality", PI=FALSE) +  
  autolayer(SARIMA_forecast, series="SARIMA Forecasting", PI=FALSE) +  
  autolayer(SSES_original, series="Exponential Smooth Forecasting", PI=FALSE) +  
  ylab("Hydro flow, m3/s") +  
  xlab("Year") +  
  labs(col="Types")
```

```
## Warning in ggplot2::geom_line(ggplot2::aes_(group = ~series, colour =  
## ~series), : Ignoring unknown parameters: `PI`  
  
## Warning in ggplot2::geom_line(ggplot2::aes_(x = ~timeVal, y = ~seriesVal, :  
## Ignoring unknown parameters: `PI`
```



Q9

From the plot in Q9 which model or model(s) are leading to the better forecasts? Explain your answer. Hint: Think about which models are doing a better job forecasting the high and low inflow months for example.

Answer: SARIMA model is leading to the better forecast because it did a better job forecasting the high and low inflow months. The difference between the SARIMA forecast's values and the actual historical data values is the smallest compared to the other models (the highest point and lowest point in the graph).

Q10

Now compute the following forecast metrics we learned in class: RMSE and MAPE, for all the models you plotted in part Q9. You can do this by hand since you have forecasted and observed values for the year of 2010. Or you can use R function `accuracy()` from package "forecast" to do it. Build a table with the results and highlight the model with the lowest MAPE. Does the lowest MAPE corresponds match your answer for part Q10?

```
last_obs1 <- as.numeric(inflow_data_only_2010[,2])

#Model 1: ARIMA Forecasting with seasonality
ARIMA_forecast_scores <- accuracy(Sum_Sea_forecast,last_obs1) #store the performance metrics

#Model 2: SARIMA forecast
SARIMA_forecast_scores <- accuracy(SARIMA_forecast$mean,last_obs1)

# Model 3: SES
SES_scores <- accuracy(SES_original$mean,last_obs1)
```

Table 1: Forecast Accuracy for Seasonal Data

	ME	RMSE	MAE	MPE	MAPE
ARIMA	-5818.450	5852.520	5818.450	-38.04169	38.04169
SARIMA	-4031.818	4171.259	4031.818	-23.87315	23.87315
SES	-2165.000	11888.481	10718.167	-59.74209	83.65282

```
#create data frame
seas_scores <- as.data.frame(rbind(ARIMA_forecast_scores, SARIMA_forecast_scores, SES_scores))
row.names(seas_scores) <- c("ARIMA", "SARIMA", "SES")

kbl(seas_scores,
    caption = "Forecast Accuracy for Seasonal Data",
    digits = array(5, ncol(seas_scores))) %>%
kable_styling(full_width = FALSE, position = "center") %>%
#highlight model with lowest MAPE
kable_styling(latex_options="striped", stripe_index = which.min(seas_scores[, "MAPE"]))

#choose model with lowest MAPE
best_model_index <- which.min(seas_scores[, "MAPE"])
cat("The best model by MAPE is:", row.names(seas_scores[best_model_index,]))
```

The best model by MAPE is: SARIMA

Answer: Yes, the lowest MAPE corresponds match my answer for part Q9 that SARIMA model is the best models by MAPE.